Revenue Intelligence System

I built a Snowflake-native Revenue Intelligence application that models RAW operational data into curated MARTS fact tables, defines governed GTM metrics (ARR from MRR, cohort-based NRR/GRR with boundary-month safeguards, win rate, and pipeline coverage), and embeds lightweight data quality checks to ensure interpretability. On top of this foundation, I implemented a reusable semantic layer and validated lineage and dependencies in Snowsight to reinforce metric governance. I then integrated Snowflake Cortex (AI_COMPLETE) to generate a board-ready executive narrative and a strict, evidence-grounded Analyst Q&A experience with confidence scoring, ensuring AI outputs remain fully tied to governed data. The result is an end-to-end, AI-ready analytics system that demonstrates semantic modeling, observability, and responsible AI usage directly within the Snowflake platform.

Github Link: https://github.com/bhoomikaa/gtm-revenue-copilot

Deployed a Streamlit app inside Snowflake for exec + GTM analytics.

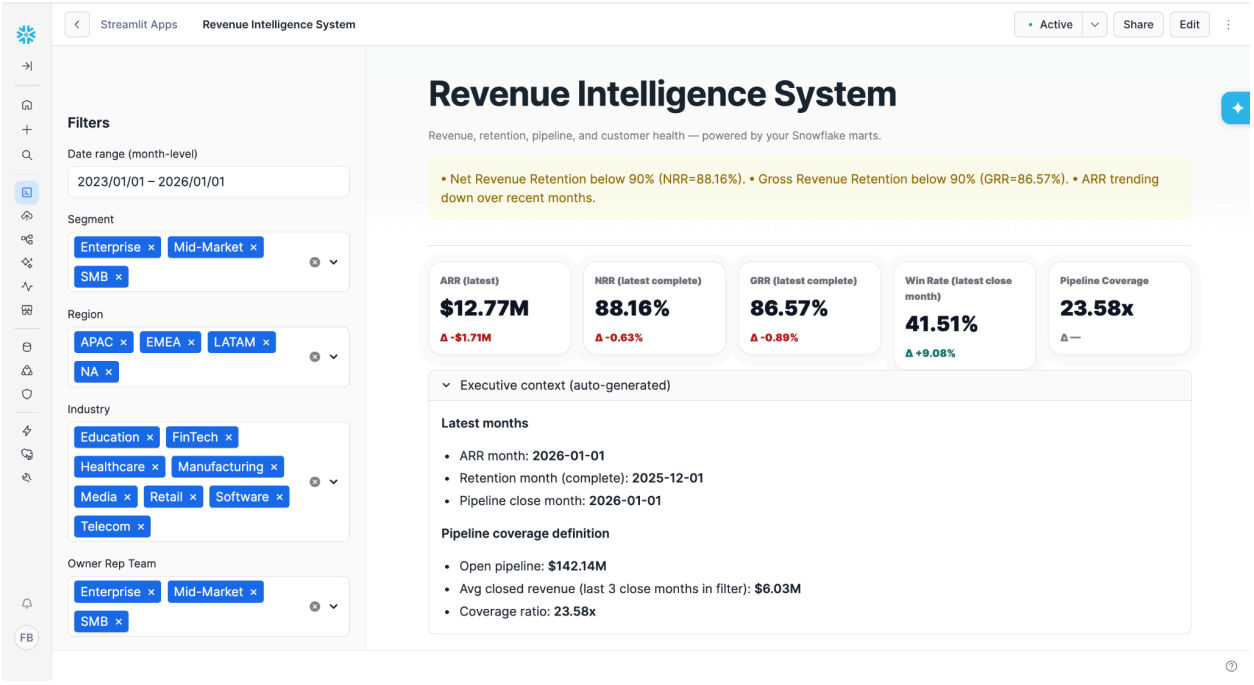

FEATURES:

Board-ready executive summary combining governed ARR, retention, win rate, and pipeline coverage metrics with Cortex-generated narrative grounded in Snowflake MARTS data.

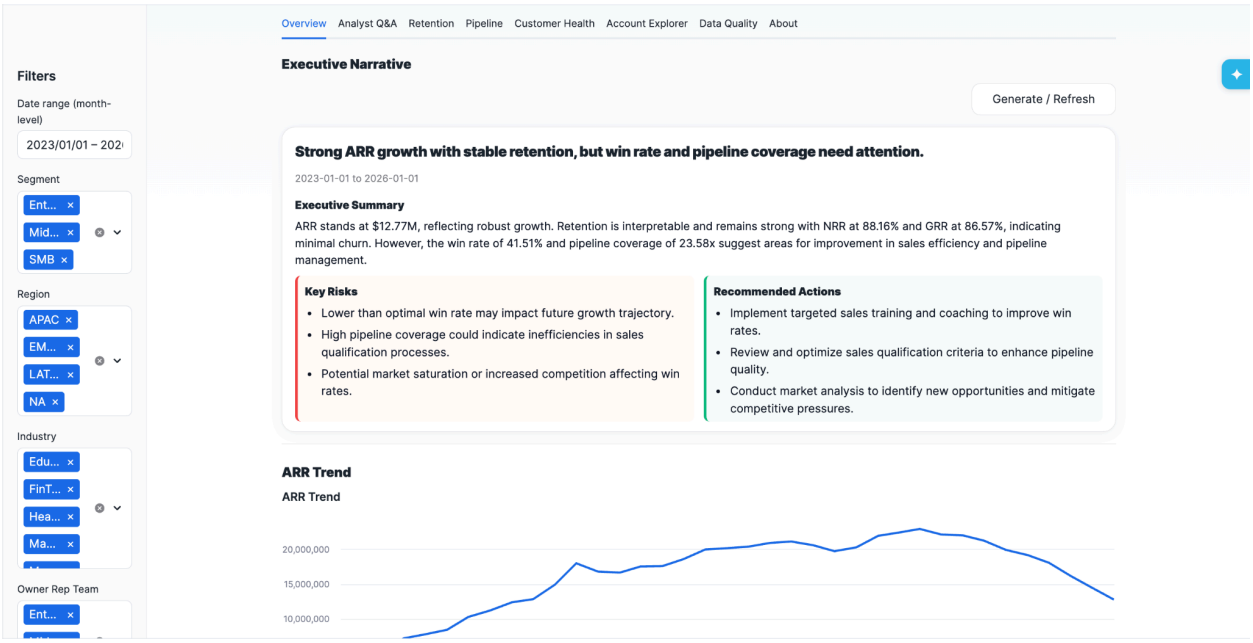

Cortex-powered, strictly grounded analytics assistant that answers GTM questions using a governed JSON data pack with evidence bullets and confidence scoring.

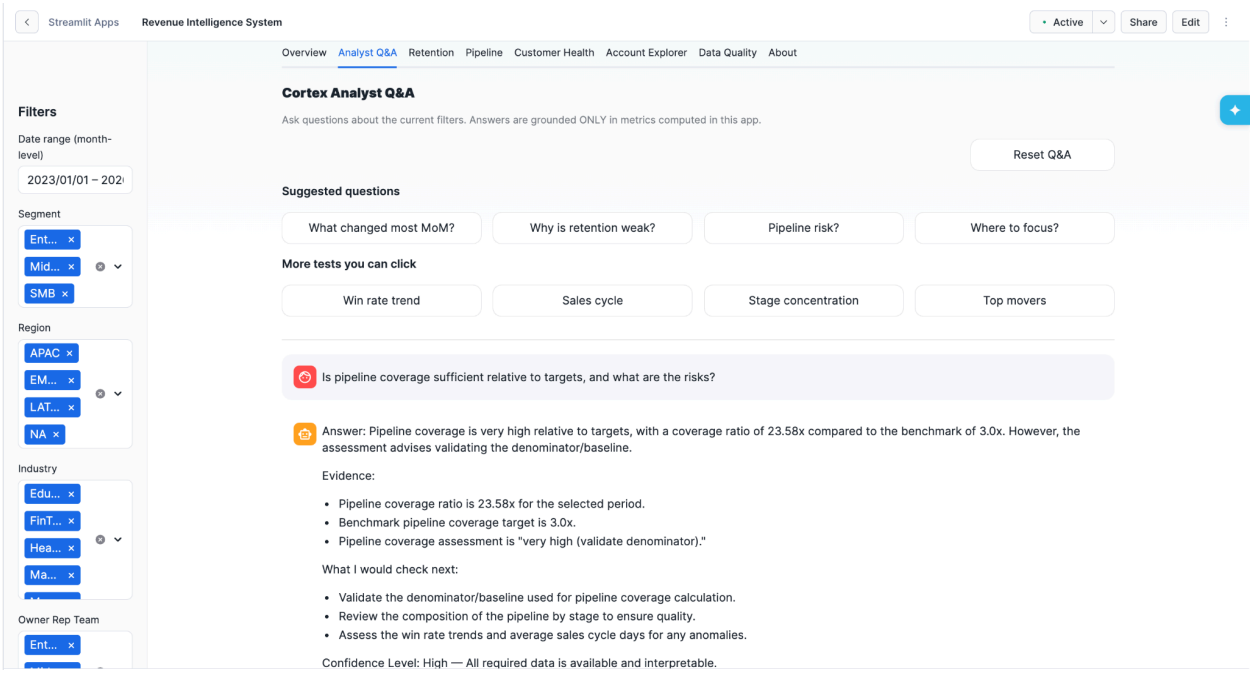

Cohort-based NRR/GRR analysis with boundary-month safeguards, plus MRR movement breakdown (expansion, contraction, churn) built directly from curated fact tables.

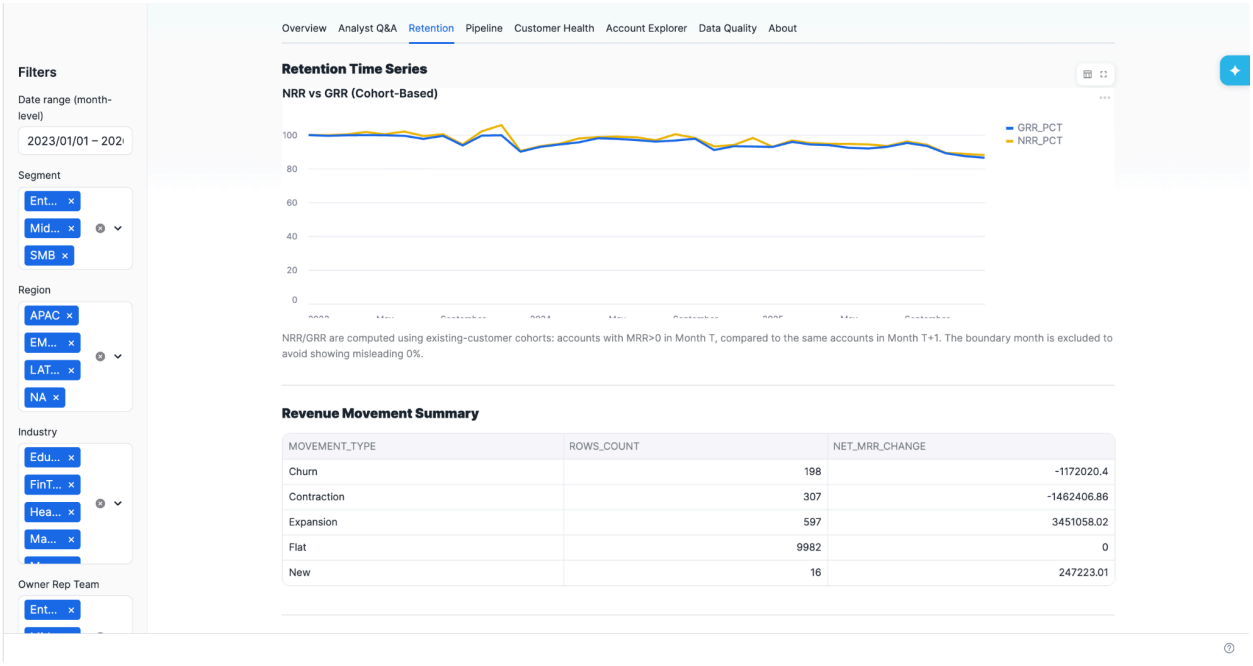

MOVEMENT_TYPE	ROWS_COUNT	NET_MRR_CHANGE
Churn	198	-1172020.4
Contraction	307	-1462406.86
Expansion	597	3451058.02
Flat	9982	0
New	16	247223.01

Lightweight observability layer showing date bounds, row counts, and sanity checks to validate data completeness and metric reliability.

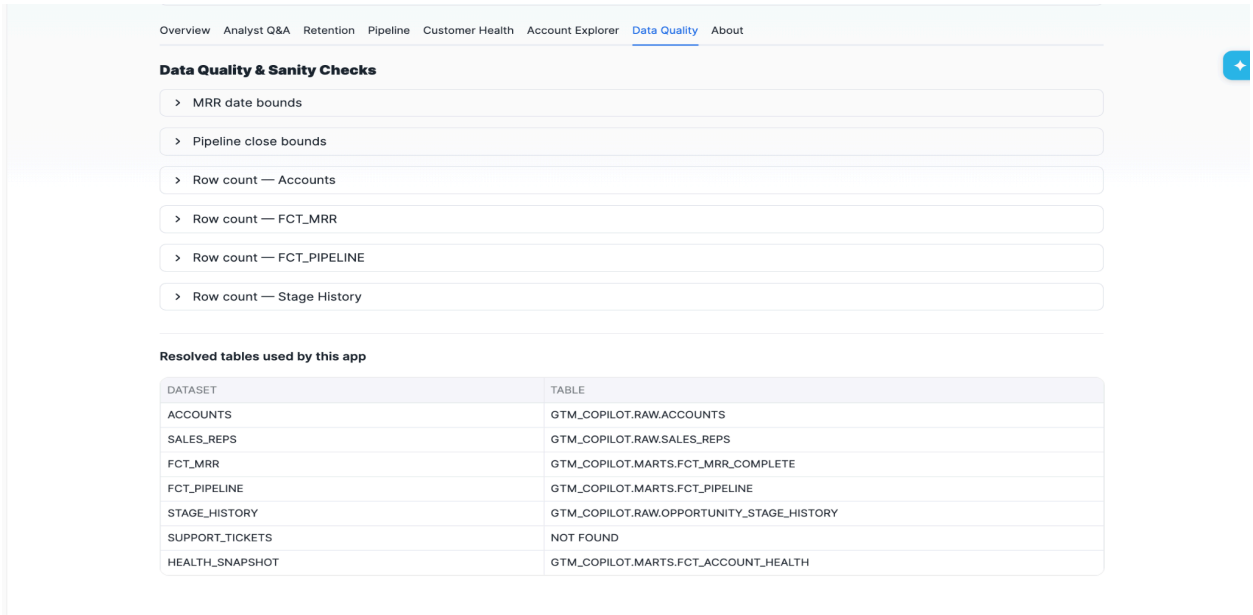

DATASET	TABLE
ACCOUNTS	GTM_COPILOT.RAW.ACCOUNTS
SALES_REPS	GTM_COPILOT.RAW.SALES_REPS
FCT_MRR	GTM_COPILOT.MARTS.FCT_MRR_COMPLETE
FCT_PIPELINE	GTM_COPILOT.MARTS.FCT_PIPELINE
STAGE_HISTORY	GTM_COPILOT.RAW.OPPORTUNITY_STAGE_HISTORY
SUPPORT_TICKETS	NOT FOUND
HEALTH_SNAPSHOT	GTM_COPILOT.MARTS.FCT_ACCOUNT_HEALTH

FLOW:

1)Running analytics workloads with RBAC and warehouse-managed compute in Snowsight.

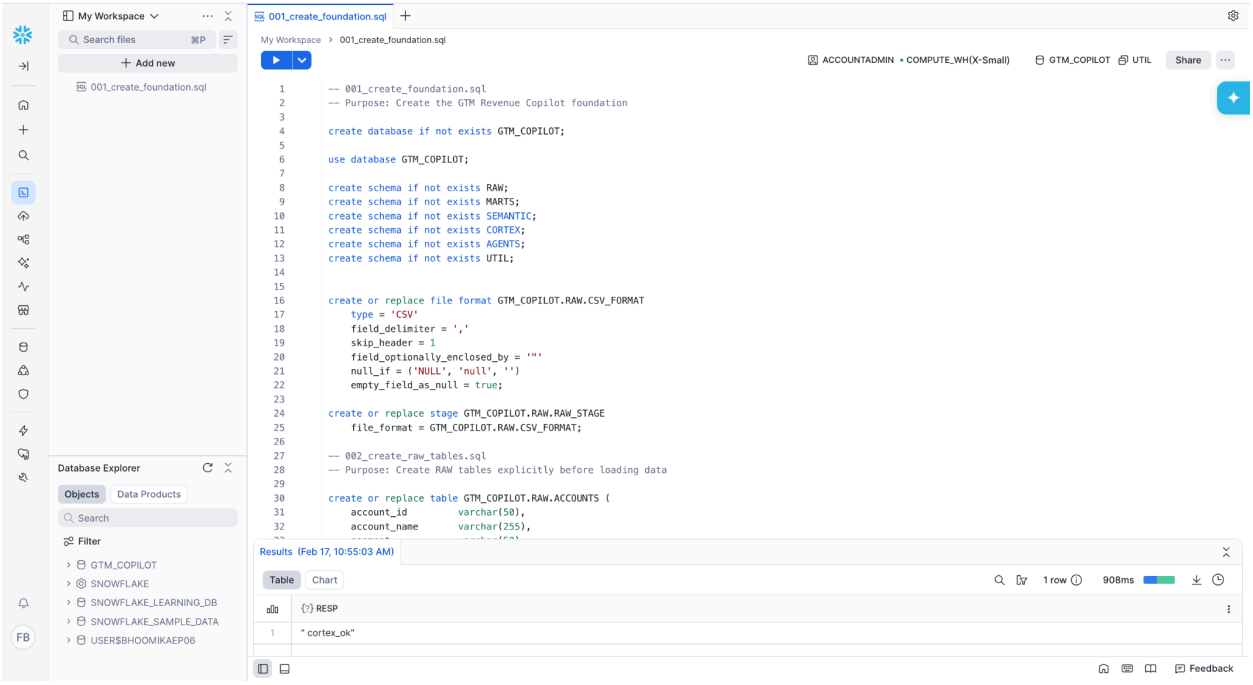

2)Warehouse layering: RAW ingestion → MARTS curated facts for consistent metrics.

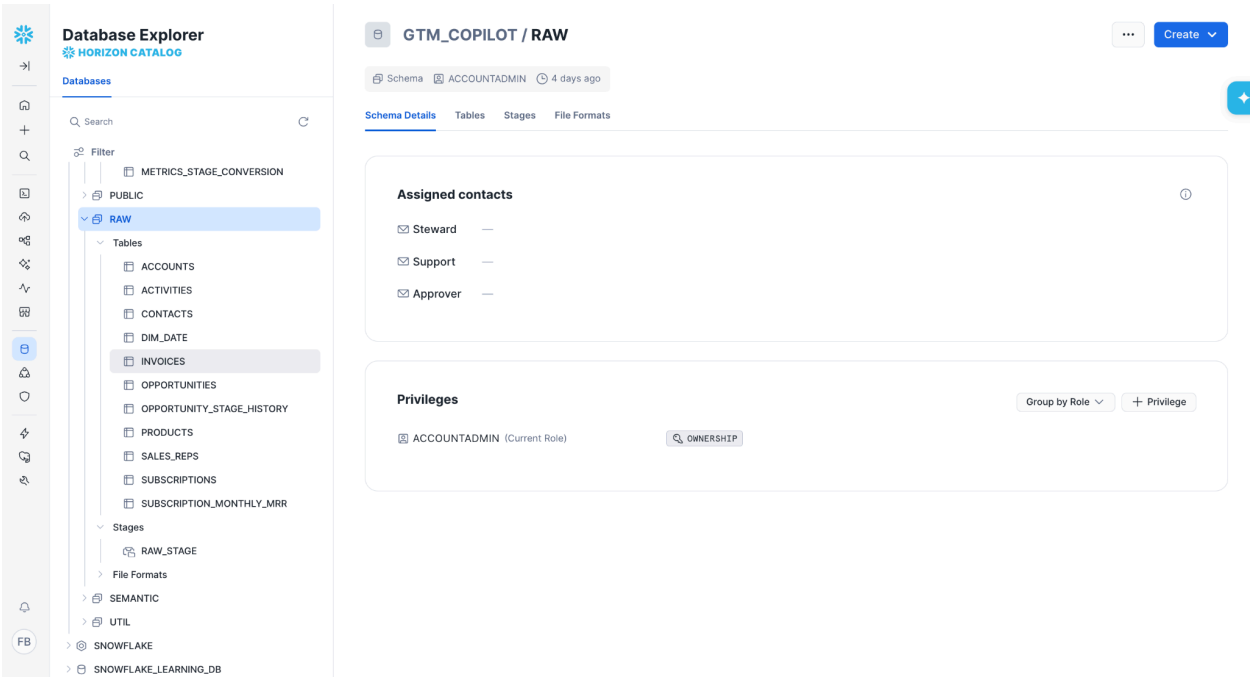

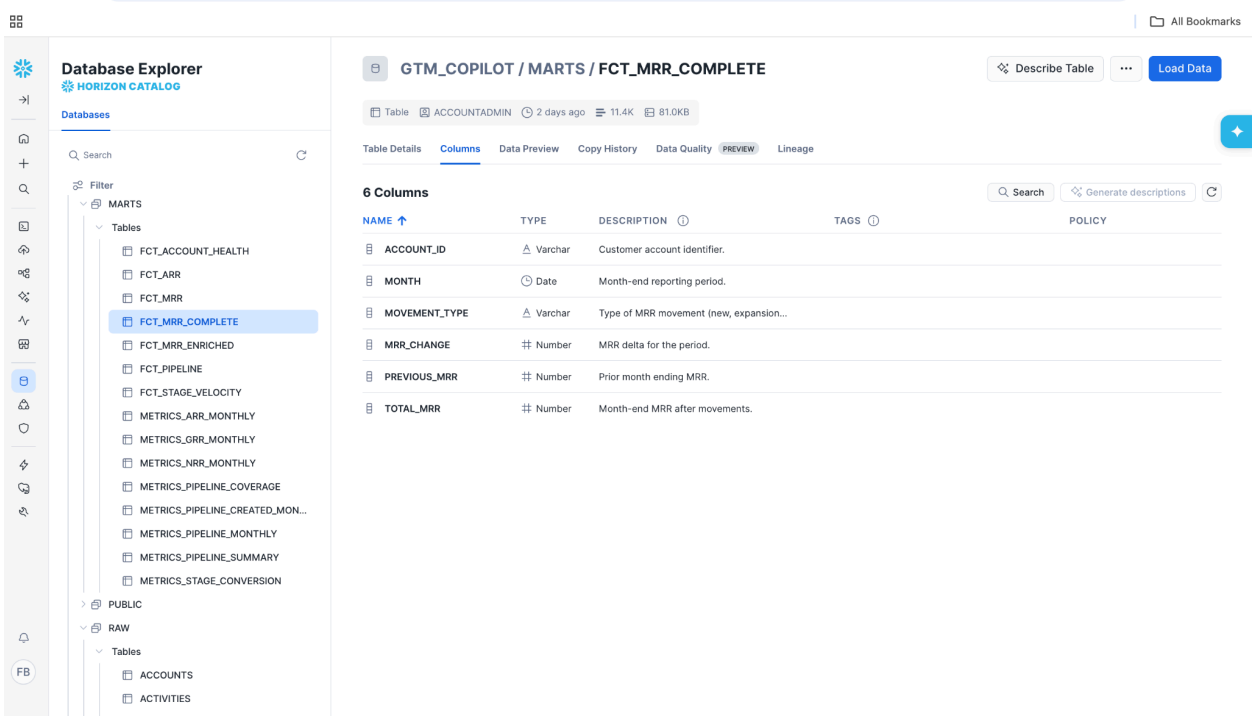

3)Curated fact table with stable, metric-ready fields (MRR by account and month).

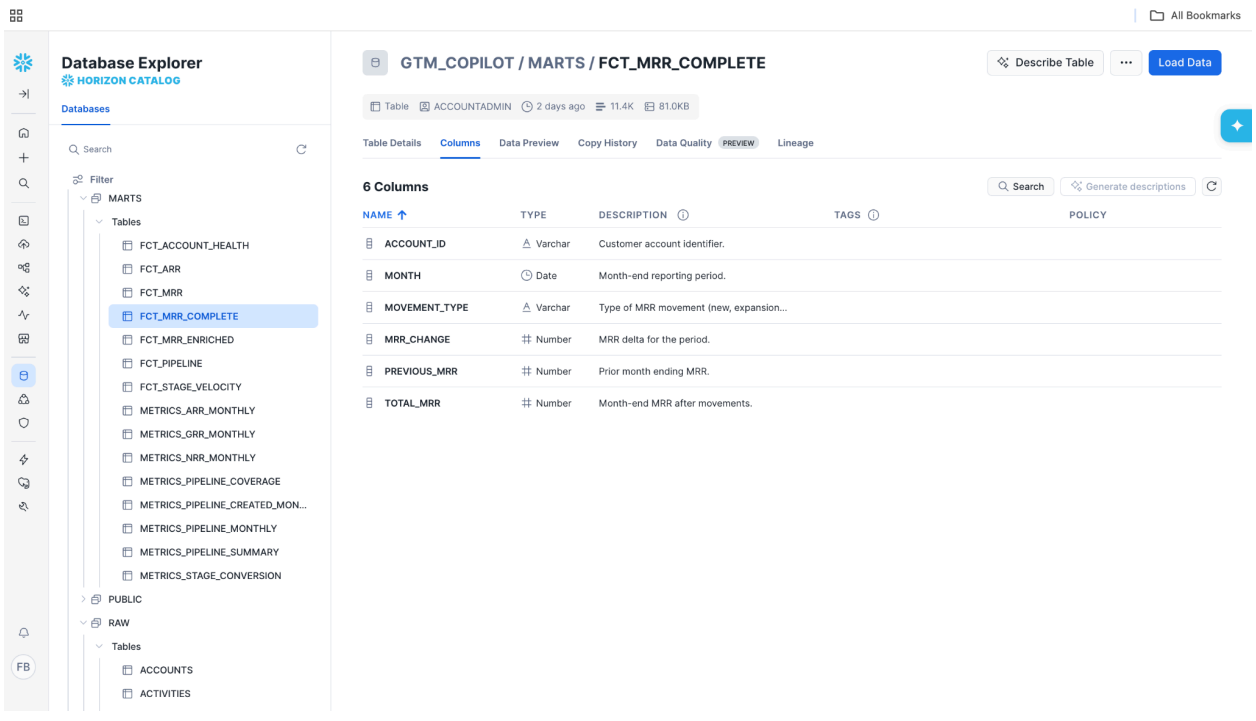

4)Pipeline fact table supporting win rate, sales cycle, and coverage analysis.

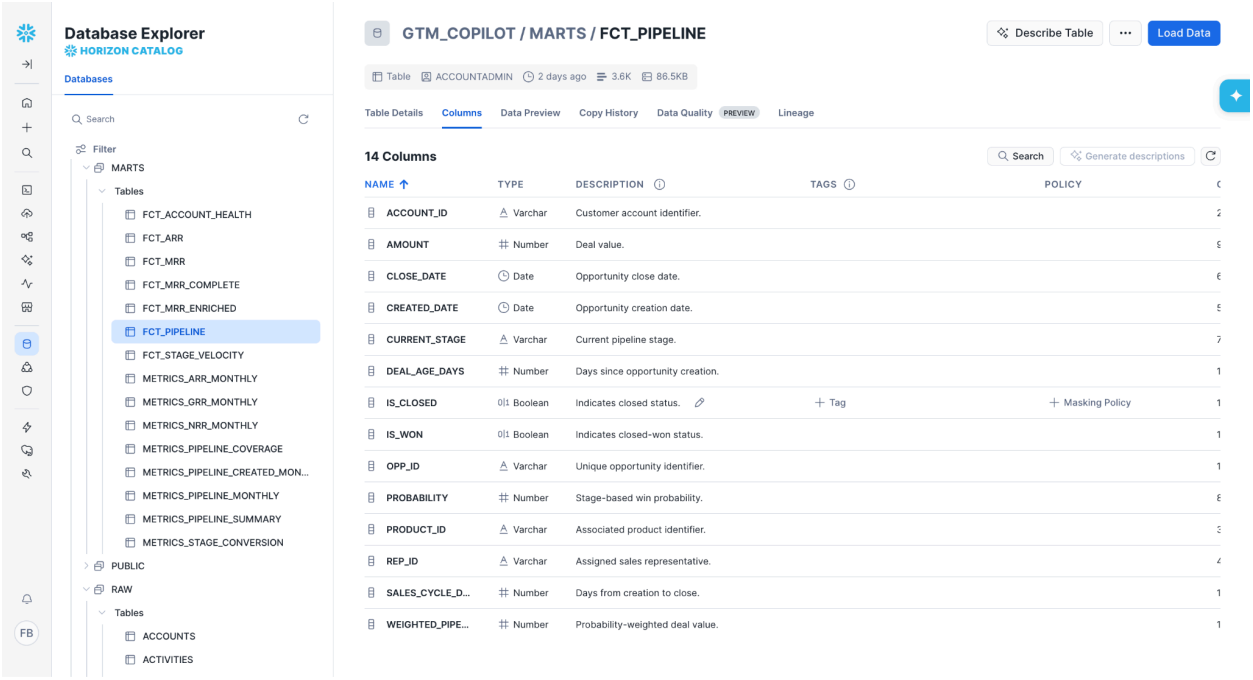

5)Cohort-correct NRR/GRR logic using month T → month T+1 joins to prevent boundary-month distortion.

```sql
-- 104_metrics_nrr_monthly.sql
-- Purpose: Monthly Net Revenue Retention (NRR) across existing customer base and excluding incomplete months
create or replace table GTM_COPILOT.MARTS.METRICS_NRR_MONTHLY as

with max_valid_month as (
    select max(month) as max_month
    from GTM_COPILOT.MARTS.FCT_MRR_COMPLETE
    where total_mrr > 0
),
base as (
    select
        account_id,
        month,
        total_mrr,
        lag(total_mrr) over (partition by account_id order by month) as previous_mrr
    from GTM_COPILOT.MARTS.FCT_MRR_COMPLETE
),
eligible as (
    select
        account_id,
        month,
        total_mrr,
        previous_mrr
    from base
    where previous_mrr > 0
)

select
    e.month,
    sum(e.previous_mrr) as start_mrr,
    sum(e.total_mrr) as end_mrr,
    round((sum(e.total_mrr) / sum(e.previous_mrr)) * 100, 2) as nrr_pct
from eligible e
join max_valid_month m
    on e.month <= m.max_month
group by e.month
order by e.month;
```

6)Validated monthly NRR outputs computed directly from curated MARTS facts.

```
428
429  select *
430  from GTM_COPILOT.MARTS.METRICS_NRR_MONTHLY
431  order by month desc
432  limit 5;
433
```

Results (just now)

Table | Chart — 5 rows — 75ms

	MONTH	START_MRR	END_MRR	NRR_PCT
1	2026-01-01	1206672.41	1063853.77	88.16
2	2025-12-01	1347062.10	1196043.16	88.79
3	2025-11-01	1506320.08	1347062.10	89.43
4	2025-10-01	1598047.84	1506320.08	94.26
5	2025-09-01	1660803.53	1597417.48	96.18

7)Query Profile validating scan nodes, execution plan, and warehouse compute efficiency.

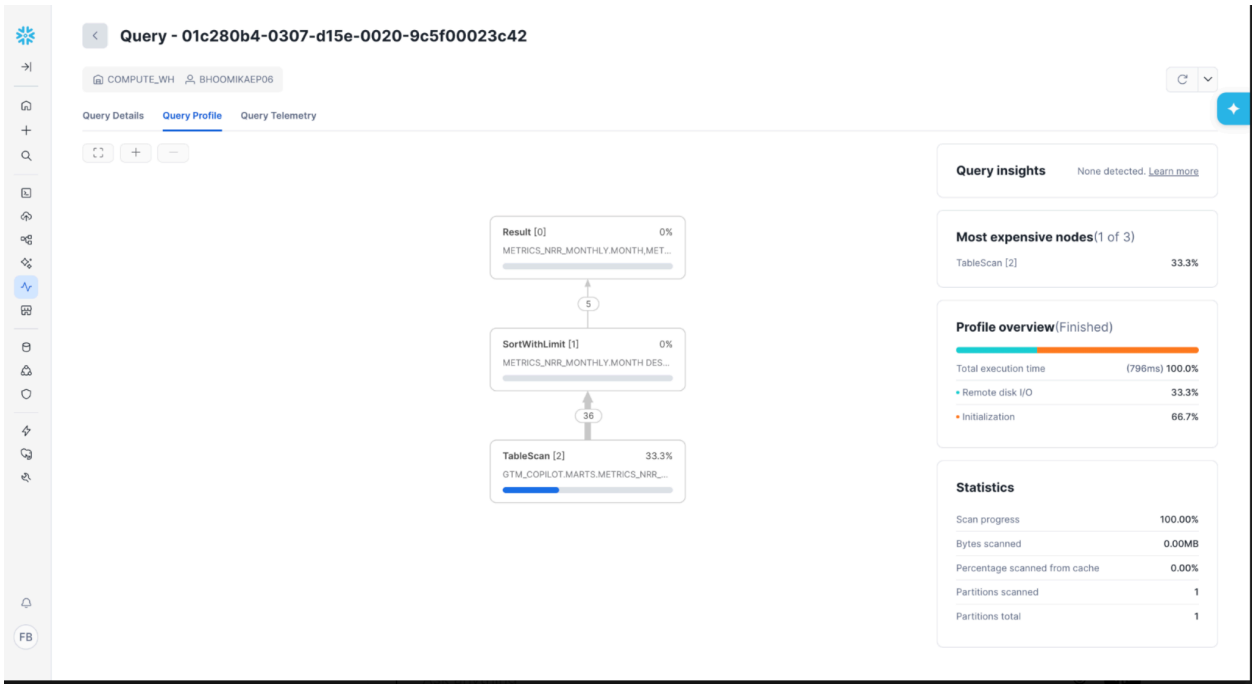

8)Cortex LLM execution (mistral-large2) embedded in Snowflake for contextual analytics reasoning.

```
select AI_COMPLETE(
  'mistral-large2',
  'Explain what Net Revenue Retention (NRR) means in one sentence.'
) as resp;
```